

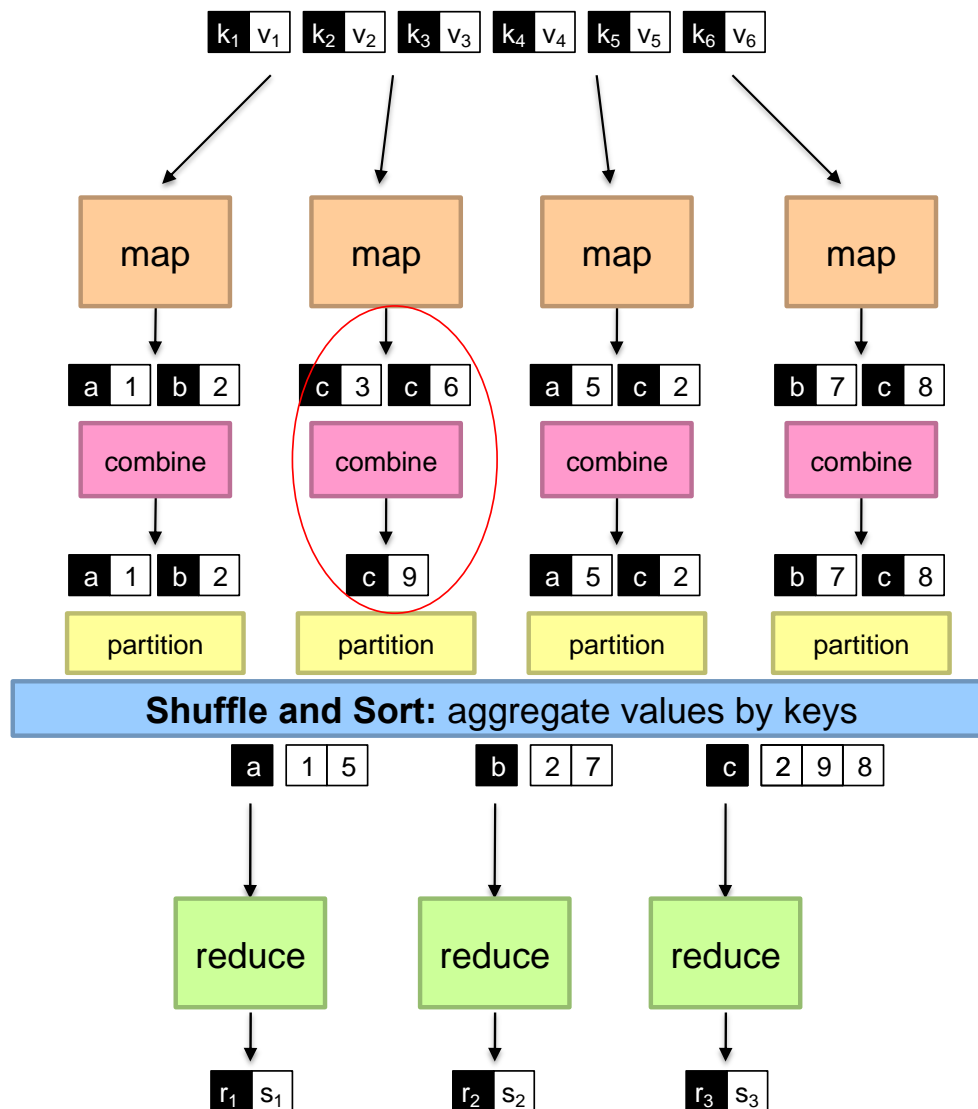
# Tutorial 3: MapReduce

Li Yuan

*li.yuan@u.nus.edu*



# Review



# Question 1

- Suppose our input data to a map-reduce operation consists of integer values (the keys are not important).
- The map function takes an integer  $i$  and produces the list of pairs  $(p,i)$  such that  $p$  is a prime divisor of  $i$ . For example,  $\text{map}(12) = [(2,12), (3,12)]$ .
- The reduce function is addition. That is reduce  $(p, [i_1, i_2, \dots, i_k])$  is  $(p, i_1 + i_2 + \dots + i_k)$ .
- Compute the output, if the input is the set of integers 15, 21, 24, 30, 49. Then, identify all the pairs in the output.

# Solution 1

- *Map does the following:*
  - 15 -> (3,15), (5,15)
  - 21 -> (3,21), (7,21)
  - 24 -> (2,24), (3,24)
  - 30 -> (2,30), (3,30), (5,30)
  - 49 -> (7,49)
- *Then group by keys:*
  - (2, [24,30])
  - (3, [15,21,24,30])
  - (5, [15,30])
  - (7, [21,49])
- *Reduce add elements:*
  - (2,54)
  - (3,90)
  - (5,45)
  - (7,70)

## Question 2

- Using the matrix-vector multiplication described in Section 2.3.1, applied to the matrix and vector:

$$\begin{bmatrix} 1 & 2 & 3 & 4 \\ 5 & 6 & 7 & 8 \\ 9 & 10 & 11 & 12 \\ 13 & 14 & 15 & 16 \end{bmatrix} \begin{bmatrix} 1 \\ 2 \\ 3 \\ 4 \end{bmatrix}$$

- Apply the Map function to this matrix and vector.
- Then, identify all the key-value pairs that are output of Map.

# Solution 2

## *Section 2.3.1: Matrix-Vector Multiplication by MapReduce*

- Suppose we have an  $n \times n$  matrix  $M$ , whose element in row  $i$  and column  $j$  will be denoted  $m_{ij}$ . Suppose we also have a vector  $v$  of length  $n$ , whose  $j$ th element is  $v_j$ . Then the matrix-vector product is the vector  $x$  of length  $n$ , whose  $i$ th element  $x_i$  is given by  $x_i = \sum_{j=1}^n m_{ij} \cdot v_j$ .
- The Map Function: for each matrix element  $m_{ij}$ , produce the key-value pair  $(i, m_{ij} \cdot v_j)$
- The Reduce Function: sum all the values associated with a given key  $i$ , and the result is a pair  $(i, x_i)$

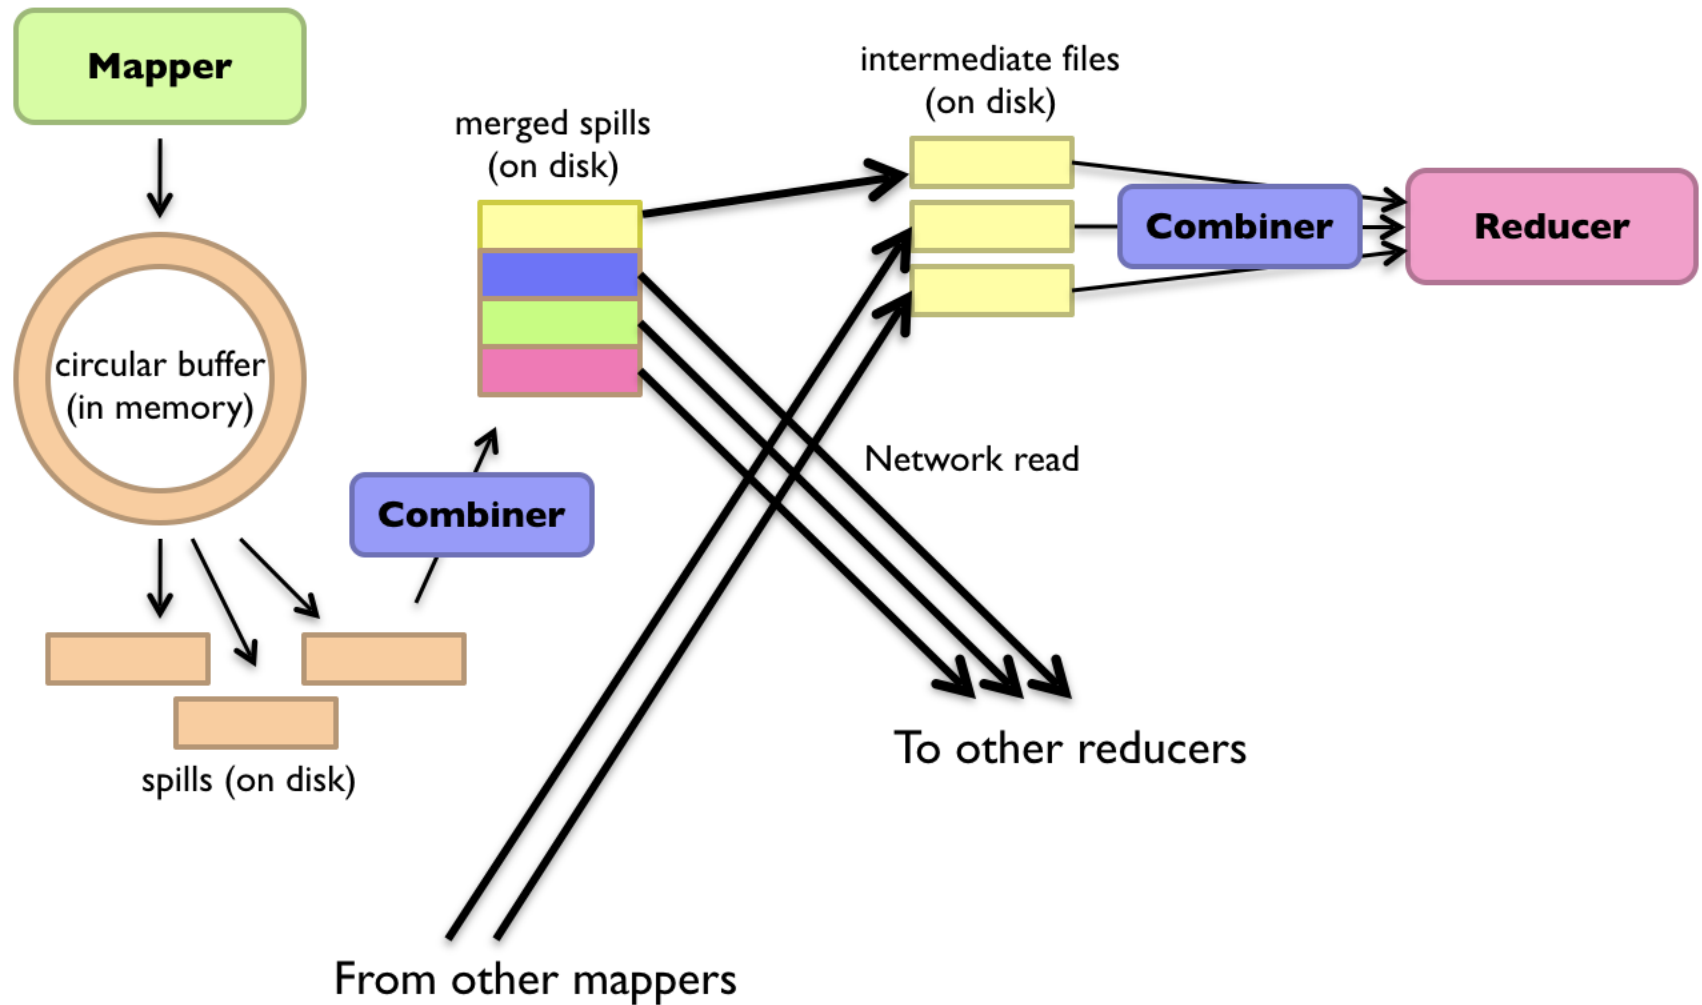
## Solution 2

- Thus, in row-major order, the sixteen key-value pairs produced are:

$$\begin{bmatrix} (1,1) & (1,4) & (1,9) & (1,16) \\ (2,5) & (2,12) & (2,21) & (2,32) \\ (3,9) & (3,20) & (3,33) & (3,48) \\ (4,13) & (4,28) & (4,45) & (4,64) \end{bmatrix}$$

- Reduce will sum the values corresponding to the same key, thus, results of reduce are: (1,30), (2,70), (3,110), (4,150)

# Shuffle and Sort





## Question 3

- **Consider a simple example:**
  - We have a large dataset where input keys are strings and input values are integers.
  - We wish to compute the mean of all integers associated with the same key (rounded to the nearest integer).
  - A real-world example might be a large user log from a popular website, where keys represent user ids and values represent some measure of activity such as elapsed time for a particular session.
  - A program Tommy has implemented the problem on MapReduce. He has written a few versions with the pseudo code shown in Figures 1—4.

## Question 3a

```
1: class MAPPER
2:   method MAP(string  $t$ , integer  $r$ )
3:     EMIT(string  $t$ , integer  $r$ )
1: class REDUCER
2:   method REDUCE(string  $t$ , integers  $[r_1, r_2, \dots]$ )
3:      $sum \leftarrow 0$ 
4:      $cnt \leftarrow 0$ 
5:     for all integer  $r \in$  integers  $[r_1, r_2, \dots]$  do
6:        $sum \leftarrow sum + r$ 
7:        $cnt \leftarrow cnt + 1$ 
8:      $r_{avg} \leftarrow sum / cnt$ 
9:     EMIT(string  $t$ , integer  $r_{avg}$ )
```

**a) Initially, Tommy has finished an implementation with Version 1 (Figure 1). He finds that the implementation can have correct results, but the performance is very slow. Why?**

## Solution 3a

**It requires shuffling all key-value pairs from mappers to reducers across the network, which is highly inefficient.**

## Question 3b

```
1: class MAPPER
2:   method MAP(string  $t$ , integer  $r$ )
3:     EMIT(string  $t$ , integer  $r$ )
1: class COMBINER
2:   method COMBINE(string  $t$ , integers  $[r_1, r_2, \dots]$ )
3:      $sum \leftarrow 0$ 
4:      $cnt \leftarrow 0$ 
5:     for all integer  $r \in$  integers  $[r_1, r_2, \dots]$  do
6:        $sum \leftarrow sum + r$ 
7:        $cnt \leftarrow cnt + 1$ 
8:     EMIT(string  $t$ , pair ( $sum, cnt$ ))           ▷ Separate sum and count
1: class REDUCER
2:   method REDUCE(string  $t$ , pairs  $[(s_1, c_1), (s_2, c_2) \dots]$ )
3:      $sum \leftarrow 0$ 
4:      $cnt \leftarrow 0$ 
5:     for all pair  $(s, c) \in$  pairs  $[(s_1, c_1), (s_2, c_2) \dots]$  do
6:        $sum \leftarrow sum + s$ 
7:        $cnt \leftarrow cnt + c$ 
8:      $r_{avg} \leftarrow sum / cnt$ 
9:     EMIT(string  $t$ , integer  $r_{avg}$ )
```

**b) Tommy wants to improve the performance using combiner. He comes out the second implementation (Version 2 in Figure 2). He finds that he can seldom get the reasonable results. Why?**

## Solution 3b

- Combiners are optimizations that cannot change the correctness of the algorithm.
- Combiner must have the same input and output key-value type
- If Combiner removed, the output value type of the mapper is integer, so the reducer expects to receive a list of integers as values. But the **reducer actually expects a list of pairs!**
- The correctness of the algorithm is contingent on the combiner running on the output of the mappers, and more specifically, that the combiner is run exactly once.

## Question 3c

```
1: class MAPPER
2:   method MAP(string  $t$ , integer  $r$ )
3:     EMIT(string  $t$ , pair ( $r$ , 1))

1: class COMBINER
2:   method COMBINE(string  $t$ , pairs  $[(s_1, c_1), (s_2, c_2) \dots]$ )
3:      $sum \leftarrow 0$ 
4:      $cnt \leftarrow 0$ 
5:     for all pair  $(s, c) \in$  pairs  $[(s_1, c_1), (s_2, c_2) \dots]$  do
6:        $sum \leftarrow sum + s$ 
7:        $cnt \leftarrow cnt + c$ 
8:     EMIT(string  $t$ , pair ( $sum$ ,  $cnt$ ))

1: class REDUCER
2:   method REDUCE(string  $t$ , pairs  $[(s_1, c_1), (s_2, c_2) \dots]$ )
3:      $sum \leftarrow 0$ 
4:      $cnt \leftarrow 0$ 
5:     for all pair  $(s, c) \in$  pairs  $[(s_1, c_1), (s_2, c_2) \dots]$  do
6:        $sum \leftarrow sum + s$ 
7:        $cnt \leftarrow cnt + c$ 
8:      $r_{avg} \leftarrow sum / cnt$ 
9:     EMIT(string  $t$ , pair ( $r_{avg}$ ,  $cnt$ ))
```

**c) After careful design, Tommy finally develops an efficient and correct implementation (Version 3 in Figure 3). Analyze the correctness of the combiner and efficiency of the algorithm (i.e., why it is more efficient than Version 1).**

## Solution 3c

- In the mapper we emit as the value a pair consisting of the integer and one—this corresponds to a partial count over one instance.
- The combiner separately aggregates the partial sums and the partial counts (as before), and emits pairs with updated sums and counts.
- The reducer is similar to the combiner, except that the mean is computed at the end.
- In essence, this algorithm transforms a non-associative operation (mean of numbers) into an associative operation (element-wise sum of a pair of numbers, with an additional division at the very end).

## Question 3d

```
1: class MAPPER
2:   method INITIALIZE
3:      $S \leftarrow \text{new ASSOCIATIVEARRAY}$ 
4:      $C \leftarrow \text{new ASSOCIATIVEARRAY}$ 
5:   method MAP(string  $t$ , integer  $r$ )
6:      $S\{t\} \leftarrow S\{t\} + r$ 
7:      $C\{t\} \leftarrow C\{t\} + 1$ 
8:   method CLOSE
9:     for all term  $t \in S$  do
10:        $\text{EMIT}(\text{term } t, \text{pair } (S\{t\}, C\{t\}))$ 
```

**d) Tommy analyzes the efficiency of Version 3, and comes out an even more efficient implementation (Version 4 in Figure 4). Why does Version 4 is even more efficient than Version 3?**



## Solution 3d

- Inside the mapper, the partial sums and counts associated with each string are held in memory across input key-value pairs.
- Intermediate key-value pairs are emitted only after the entire input split has been processed; similar to before, the value is a pair consisting of the sum and count.

# Summary

- **Mapper and Reducer is the key operation of divide and conquer.**
- **Combiner is an optimization step to reduce the amount of data transmission.**
- **Leverage the design of MapReduce program (How to partition the data).**

# Acknowledgement



Thanks to Li Qinbin for making these slides.

*liqinbin@u.nus.edu*