

Tutorial 8: Large Graph Processing II


Li Yuan

li.yuan@u.nus.edu



Problem 1

1. Given the following graph,
 - 1) how many dead ends are there in the graph? For each dead end (if any), please indicate the set of vertices forming the dead end.
 - 2) how many spider traps are there in the graph? For each spider trap (if any), please indicate the set of vertices forming the spider trap.



2 problems:

(1) Some pages are
dead ends (have no out-links)


Random walk has “nowhere” to go to

Such pages cause importance to “leak out”

(2) **Spider traps:**
(all out-links are within the group)


Random walked gets “stuck” in a trap

And eventually spider traps absorb all importance




Solution 1


1) No dead ends



Solution 1

- 1) No dead ends
- 2) Two spider traps






$$A = \beta M + (1 - \beta) \left[\frac{1}{N} \right]_{N \times N}$$

$$\beta \times \begin{matrix} & \begin{matrix} y & a & m \end{matrix} \\ \begin{matrix} y \\ a \\ m \end{matrix} & \begin{bmatrix} \frac{1}{2} & \frac{1}{2} & 0 \\ \frac{1}{2} & 0 & 0 \\ 0 & \frac{1}{2} & 1 \end{bmatrix} \end{matrix} + (1 - \beta) \times \begin{matrix} & \begin{matrix} y & a & m \end{matrix} \\ \begin{matrix} y \\ a \\ m \end{matrix} & \begin{bmatrix} \frac{1}{3} & \frac{1}{3} & \frac{1}{3} \\ \frac{1}{3} & \frac{1}{3} & \frac{1}{3} \\ \frac{1}{3} & \frac{1}{3} & \frac{1}{3} \end{bmatrix} \end{matrix} = A$$

$$\begin{bmatrix} r_y \\ r_a \\ r_m \end{bmatrix} = A \cdot \begin{bmatrix} r_y \\ r_a \\ r_m \end{bmatrix}$$

Problem 2


Set up the PageRank equations for the below graph, assuming $\beta = 0.8$ (jump probability = $1 - \beta$). Denote the PageRank of node x by $r(x)$.



Solution 2

$$M = \begin{bmatrix} 0 & 0 & \frac{1}{3} & \frac{1}{2} \\ 1 & 0 & \frac{1}{3} & 0 \\ 0 & 1 & 0 & \frac{1}{2} \\ 0 & 0 & \frac{1}{3} & 0 \end{bmatrix}$$

$$N = \begin{bmatrix} \frac{1}{4} & \frac{1}{4} & \frac{1}{4} & \frac{1}{4} \\ \frac{1}{4} & \frac{1}{4} & \frac{1}{4} & \frac{1}{4} \\ \frac{1}{4} & \frac{1}{4} & \frac{1}{4} & \frac{1}{4} \\ \frac{1}{4} & \frac{1}{4} & \frac{1}{4} & \frac{1}{4} \end{bmatrix}$$



Solution 2

$$A = \beta M + (1 - \beta)N$$

$$r = Ar = \beta Mr + (1 - \beta)Nr$$

$$r(A) = \frac{4}{15}r(C) + \frac{2}{5}r(D) + \frac{1}{20}(r(A) + r(B) + r(C) + r(D))$$

$$r(B) = \frac{4}{5}r(A) + \frac{4}{15}r(C) + \frac{1}{20}(r(A) + r(B) + r(C) + r(D))$$

$$r(C) = \frac{4}{5}r(B) + \frac{2}{5}r(D) + \frac{1}{20}(r(A) + r(B) + r(C) + r(D))$$

$$r(D) = \frac{4}{15}r(C) + \frac{1}{20}(r(A) + r(B) + r(C) + r(D))$$

Problem 3

3. Suppose you have a large graph, and you will implement breadth first traversal on the graph. Each vertex has three attributes: 1) *id*, the vertex ID, 2) *isVisited*, indicating the vertex has been visited or not, and 3) *vList*, the list of neighbour vertices. Show the pseudo code on how you would use Pregel to perform a breadth first traversal on the graph, starting with a pre-defined vertex V_0 . You must follow the pseudo code of the below format.

```
compute (vertex v) {  
  /* your pseudo code*/  
  /* you can use two APIs given in Pregel:  
  1) getSuperStep(): Retrieves the current superstep;  
  2) voteToHalt(): After this is called, the compute() code will no  
  longer be called for this vertex.*/  
}
```

Pregel: Computational Model

Based on Bulk Synchronous Parallel (BSP)

Computational units encoded in a directed graph

Computation proceeds in a series of **supersteps**

Message passing architecture

Each vertex, at each superstep:

Receives messages directed at it from previous superstep

Executes a user-defined function (modifying state)

Emits messages to other vertices (for the next superstep)

Termination:

A vertex can choose to deactivate itself

Is “woken up” if new messages received

Computation halts when all vertices are inactive

```
compute (vertex v) {  
    if(getSuperStep() == 0){  
        if(v == V0) {  
            send messages to v's neighbors.  
            isVisited=true;  
        }  
    }  
    else if (isVisited == false) {  
        if(there is a message for v){  
            send messages to v's neighbors.  
            isVisited=true;  
        }  
    }  
    v.voteToHalt();  
}
```

Acknowledgement



Thanks to Li Qinbin for making these slides.

liqinbin@u.nus.edu